



RGBA 0 -> 0-3
1 -> 4-7
.
.
.
127 -> 508-511

we can try 16 or 32 threads/block
c.c. 20

L1 memory can be split in the following ways
16 kb - default shared memory
32 kb
48 kb
you cannot do more pixels per thread because you run
out of shared memory and it will run slowly
4096 = 128 x 32

if thread does more than 1 pixel, then write 0, 64, 128 , 192
1, 65 , 129, 193

Texture memory
- RMN story

Guidelines for choosing Blobk and Grid Design

- Guidelines:
1. The hardware should be as much occupied and busy as possible.
 - Don't waste threads with this
 2. The execution should involve as little stalling as possible
 - if else branches
 - global memory I/O
 - shared memory collisions
 - synchronisation

- Practical choises:
1. A thread should do more work that can ideally be executed in parallel by a warp (32 - threads)
 - not little kernels, rather larger
 2. Use more registers per thread to avoid using shared memory(if you can afford it) may result in a smaller number of threads per block
 - 3 The number of threads per block should be multiple of the warp size (128-256)
 4. The grid (totality of the blocks) should provide multiple(3-4) blocks per streaming of the processor(SM)
 5. Use blocks big enough to cover multiple warps

How many threads per block?

min $\left(\begin{array}{l} \text{numOfWarpSched} * \text{warpSize} \\ \text{regsPerBlock} / \text{regsPerThread} \\ \text{sharedRamSize} / \text{sharedRamPerThread} \\ \text{MaxThreadsPerSM} \end{array} \right)$ GPU model -> architecture -> specs

Streams

an[i]++; // slow if done on global memory because it actually does 1 read and 1 write

atomicAdd
atomicXchg
atomicAddAndXchg